

8. Algorithmic (automatic) differentiation

Last time

- Systems of nonlinear equations
- Derivatives in higher dimensions
- Multidimensional Newton

Goal for today

- Rules for derivatives
- Teach computer to follow rules: **algorithmic differentiation**
- Computing derivatives in higher dimensions

Derivatives as rules

- Think about differentiating

$$f(x) = \sin(x + (x + 1)(x^2 + 2))$$

Derivatives as rules

- Think about differentiating
$$f(x) = \sin(x + (x + 1)(x^2 + 2))$$
- We apply **rules** for derivatives of sums, products, compositions
- This feels like following an **algorithm**
- Can we get the computer to execute that for us?

Derivatives as rules

- Think about differentiating
$$f(x) = \sin(x + (x + 1)(x^2 + 2))$$
- We apply **rules** for derivatives of sums, products, compositions
- This feels like following an **algorithm**
- Can we get the computer to execute that for us?
- We only discuss **forward-mode algorithmic (automatic) differentiation**

Derivatives (review)

- **Derivative = rate of change**
- Increasing \iff derivative > 0

Derivatives (review)

- **Derivative = rate of change**
- Increasing \iff derivative > 0
- Tells us how function looks **locally** (close to a point)
- E.g. Use to analyze dynamics near a fixed point
- Crucial in root finding + optimization

Derivatives II

- **Derivative** of $f : \mathbb{R} \rightarrow \mathbb{R}$ at a is slope of **tangent line**
- **Tangent line**: “touches” graph of function at point
- Best affine approximation to f near a .

Derivatives II

- **Derivative** of $f : \mathbb{R} \rightarrow \mathbb{R}$ at a is slope of **tangent line**
- **Tangent line**: “touches” graph of function at point
- Best affine approximation to f near a .
- Notation: $f'(a)$, or $\left. \frac{df}{dx} \right|_a$ (if you must).

Definition of derivative

- Definition of derivative of $f : \mathbb{R} \rightarrow \mathbb{R}$ at $a \in \mathbb{R}$:

$$f'(a) := \lim_{h \rightarrow 0} \frac{f(a+h) - f(a)}{h}$$

- Intuition: Limit of slopes of **secant lines** (“rise over run”)

Definition of derivative

- Definition of derivative of $f : \mathbb{R} \rightarrow \mathbb{R}$ at $a \in \mathbb{R}$:

$$f'(a) := \lim_{h \rightarrow 0} \frac{f(a+h) - f(a)}{h}$$

- Intuition: Limit of slopes of **secant lines** (“rise over run”)
- How can we calculate derivatives numerically?

How to *calculate* derivatives numerically

- Numerically: Cannot take limit $h \rightarrow 0$

How to *calculate* derivatives numerically

- Numerically: Cannot take limit $h \rightarrow 0$
- So don't! Fix $h > 0$ to get **finite difference** approximation:

$$f'(a) \simeq \frac{f(a+h) - f(a)}{h}$$

How to *calculate* derivatives numerically

- Numerically: Cannot take limit $h \rightarrow 0$
- So don't! Fix $h > 0$ to get **finite difference** approximation:

$$f'(a) \simeq \frac{f(a+h) - f(a)}{h}$$

- How good is this approximation? See later

How to *calculate* derivatives numerically

- Numerically: Cannot take limit $h \rightarrow 0$
- So don't! Fix $h > 0$ to get **finite difference** approximation:

$$f'(a) \simeq \frac{f(a+h) - f(a)}{h}$$

- How good is this approximation? See later
- Do better: calculate derivatives exactly!

Alternative viewpoint: Little-o notation

- Rewrite by hiding annoying limit:

$$\lim_{h \rightarrow 0} \left[\frac{f(a+h) - f(a)}{h} - f'(a) \right] = 0$$

Alternative viewpoint: Little-o notation

- Rewrite by hiding annoying limit:

$$\lim_{h \rightarrow 0} \left[\frac{f(a+h) - f(a)}{h} - f'(a) \right] = 0$$

- So $f(a+h) = f(a) + hf'(a) + o(h)$

where $\frac{o(h)}{h} \rightarrow 0$ when $h \rightarrow 0$

Alternative viewpoint: Little-o notation

- Rewrite by hiding annoying limit:

$$\lim_{h \rightarrow 0} \left[\frac{f(a+h) - f(a)}{h} - f'(a) \right] = 0$$

- So $f(a+h) = f(a) + hf'(a) + o(h)$

where $\frac{o(h)}{h} \rightarrow 0$ when $h \rightarrow 0$

- $o(h)$: “some function that goes to 0 *faster* than h ”
- “Little-o notation”

Derivative from $f(a + h)$

- We showed that *if* f is differentiable at a then

$$f(a + h) = f(a) + hf'(a) + o(h)$$

Derivative from $f(a + h)$

- We showed that *if* f is differentiable at a then

$$f(a + h) = f(a) + hf'(a) + o(h)$$

- We also have the **converse**:
*Suppose $f(a + h) = A + Bh + o(h)$.
 Then $A = f(a)$ and $B = f'(a)$.*

Derivative from $f(a + h)$

- We showed that *if* f is differentiable at a then

$$f(a + h) = f(a) + hf'(a) + o(h)$$

- We also have the **converse**:
*Suppose $f(a + h) = A + Bh + o(h)$.
 Then $A = f(a)$ and $B = f'(a)$.*

- Use to calculate derivatives: calculate $f(a + h)$
- Put in this form to extract derivative as coefficient of h

Infinitesimals

- Simplify by thinking of “infinitesimal” perturbation ϵ
- With $\epsilon^2 = 0$

Infinitesimals

- Simplify by thinking of “infinitesimal” perturbation ϵ
- With $\epsilon^2 = 0$
- Manipulate first-order Taylor expansions:
- $f(a + \epsilon) = f(a) + \epsilon f'(a)$

Infinitesimals

- Simplify by thinking of “infinitesimal” perturbation ϵ
- With $\epsilon^2 = 0$
- Manipulate first-order Taylor expansions:
- $f(a + \epsilon) = f(a) + \epsilon f'(a)$
- Expand $f(a + \epsilon) = A + B\epsilon$
- Coefficient of ϵ is derivative

Sum rule for derivatives

- Let's derive some of the rules of derivatives from calculus

Sum rule for derivatives

- Let's derive some of the rules of derivatives from calculus
- Sum of two functions:

$$(f + g)(x) := f(x) + g(x)$$

Sum rule for derivatives

- Let's derive some of the rules of derivatives from calculus
- Sum of two functions:

$$(f + g)(x) := f(x) + g(x)$$

- Want $(f + g)'(a)$ so calculate $[f + g](a + \epsilon)$:

Sum rule for derivatives

- Let's derive some of the rules of derivatives from calculus
- Sum of two functions:

$$(f + g)(x) := f(x) + g(x)$$

- Want $(f + g)'(a)$ so calculate $[f + g](a + \epsilon)$:

$$\begin{aligned}[f + g](a + \epsilon) &= f(a + \epsilon) + g(a + \epsilon) \\ &= [f(a) + \epsilon f'(a)] + [g(a) + \epsilon g'(a)] \\ &= [f(a) + g(a)] + [f'(a) + g'(a)]\epsilon\end{aligned}$$

- Hence $(f + g)'(a) = (\text{coefficient of } \epsilon) = f'(a) + g'(a)$.

Product rule for derivatives

- Product of two functions:

$$(f \cdot g)(x) := f(x) \cdot g(x)$$

(Here \cdot is normal scalar multiplication)

Product rule for derivatives

- Product of two functions:

$$(f \cdot g)(x) := f(x) \cdot g(x)$$

(Here \cdot is normal scalar multiplication)

- Calculate $(f \cdot g)(a + \epsilon)$:

$$\begin{aligned} [f \cdot g](a + \epsilon) &= f(a + \epsilon) \cdot g(a + \epsilon) \\ &= [f(a) + \epsilon f'(a)] \cdot [g(a) + \epsilon g'(a)] \\ &= [f(a) \cdot g(a)] + [f(a)g'(a) + g(a)f'(a)]\epsilon. \end{aligned}$$

- Hence $(f \cdot g)'(a) = f(a)g'(a) + g(a)f'(a)$.

Derivatives from rules: Algorithmic differentiation

- For complicated function, execute each rule in turn
- E.g. $h(x) = 3x^2 + 2x = +(3 * (x * x), 2 * x)$

Derivatives from rules: Algorithmic differentiation

- For complicated function, execute each rule in turn
- E.g. $h(x) = 3x^2 + 2x = +(3 * (x * x), 2 * x)$
- What information do we need for each function?

Required information

- Fix point a where taking derivatives
- For function f , need exactly 2 pieces of information:
 - value $f(a)$
 - derivative $f'(a)$
- Represent f as

$$f \rightsquigarrow (f(a), f'(a))$$

Required information

- Fix point a where taking derivatives
- For function f , need exactly 2 pieces of information:
 - value $f(a)$
 - derivative $f'(a)$
- Represent f as

$$f \rightsquigarrow (f(a), f'(a))$$

- How can we represent this in Julia?

Representation in Julia

- Need to group together 2 pieces of information

Representation in Julia

- Need to group together 2 pieces of information
- Could use tuple or vector etc.
- But want novel **behaviour**, i.e. rules for $+$ and \times .

Representation in Julia

- Need to group together 2 pieces of information
- Could use tuple or vector etc.
- But want novel **behaviour**, i.e. rules for $+$ and \times .

- *So define a new type*
- Commonly called **dual number**

Dual number type in Julia

- Make an immutable dual number type:

Dual number type in Julia

- Make an immutable dual number type:

```
```julia
struct Dual{T<:Real}
 value::T
 deriv::T
end
```
```

. . .

- Recall: composite type is a template for box containing data
- T is a **type parameter**, any subtype of real numbers
- `Dual(a, b)` corresponds directly to $a + \epsilon b$

Creating dual numbers

- Create dual number by calling constructor, as usual:

```
julia> Dual(3, 4)  
Dual{Int64}(3, 4)
```

- Note that Julia automatically **infers** the type τ as `Int64`

Implementing arithmetic

- To implement arithmetic, import relevant functions:

```
import Base: +, *
```

- Add methods acting on objects of type `Dual`:

```
+(f::Dual, g::Dual) = Dual(f.value + g.value,  
                             f.deriv + g.deriv)
```

Differentiation using dual numbers

- Suppose have Julia function like $f(x) = x^2 + 2x$
- How can we differentiate f at $a = 3$?

Differentiation using dual numbers

- Suppose have Julia function like $f(x) = x^2 + 2x$
- How can we differentiate f at $a = 3$?
- We need ϵ components for derivative
- So want to calculate $f(a + \epsilon) = f(a) + \epsilon f'(a)$

Differentiation using dual numbers

- Suppose have Julia function like $f(x) = x^2 + 2x$
- How can we differentiate f at $a = 3$?
- We need ϵ components for derivative
- So want to calculate $f(a + \epsilon) = f(a) + \epsilon f'(a)$

Differentiation II

- Recall that $a + b\epsilon$ corresponds to $\text{Dual}(a, b)$
- So define $xx = \text{Dual}(a, 1)$
- And call $f(xx)$
- [Represents identity function $x \mapsto x$ at $x = a$, with derivative 1].

Higher-dimensional functions

- How calculate derivatives of higher-dimensional functions?
- Consider $f(x, y)$
- Suppose take dual numbers with *same* ϵ :
- Set $x = a + c\epsilon$ and $y = b + d\epsilon$
- Calculate

$$f(a + c\epsilon, b + d\epsilon)$$

Directional derivative

- For $f : \mathbb{R}^n \rightarrow \mathbb{R}$, derivative is

$$\lim_{\epsilon \rightarrow 0} \frac{1}{\epsilon} f(\mathbf{a} + \epsilon \mathbf{v})$$

- **Directional derivative** in direction \mathbf{v}

Directional derivative

- For $f : \mathbb{R}^n \rightarrow \mathbb{R}$, derivative is

$$\lim_{\epsilon \rightarrow 0} \frac{1}{\epsilon} f(\mathbf{a} + \epsilon \mathbf{v})$$

- **Directional derivative** in direction \mathbf{v}
- In 2D, define

$$g(\epsilon) := f(a + v_1 \epsilon, b + v_2 \epsilon)$$

Directional derivative II

- Expand $g(\epsilon) := f(a + v_1 \epsilon, b + v_2 \epsilon)$:

Directional derivative II

- Expand $g(\epsilon) := f(a + v_1 \epsilon, b + v_2 \epsilon)$:

$$g(\epsilon) = f(a, b) + v_1 \epsilon \frac{\partial f}{\partial x}(a, b) + v_2 \epsilon \frac{\partial f}{\partial y}(a, b)$$

- Coefficient of ϵ is

$$v_1 \frac{\partial f}{\partial x}(a, b) + v_2 \frac{\partial f}{\partial y}(a, b) = \nabla f(a, b) \cdot \mathbf{v}$$

Directional derivative II

- Expand $g(\epsilon) := f(a + v_1 \epsilon, b + v_2 \epsilon)$:

$$g(\epsilon) = f(a, b) + v_1 \epsilon \frac{\partial f}{\partial x}(a, b) + v_2 \epsilon \frac{\partial f}{\partial y}(a, b)$$

- Coefficient of ϵ is

$$v_1 \frac{\partial f}{\partial x}(a, b) + v_2 \frac{\partial f}{\partial y}(a, b) = \nabla f(a, b) \cdot \mathbf{v}$$

- How can we *calculate* this?

Calculating directional derivatives

- $f(\text{Dual}(a, v_1), \text{Dual}(b, v_2))$ calculates $\nabla f(a, b) \cdot \mathbf{v}$!

Calculating directional derivatives

- $f(\text{Dual}(a, v_1), \text{Dual}(b, v_2))$ calculates $\nabla f(a, b) \cdot \mathbf{v}$!
- $\mathbf{v} = (1, 0)$ gives $\partial f / \partial x$
- $\mathbf{v} = (0, 1)$ gives $\partial f / \partial y$

Jacobian

- For $f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$, have $f = (f_1, f_2)$ with

$$f_i(\mathbf{a} + \epsilon \mathbf{v}) = f_i(\mathbf{a}) + \epsilon \nabla f_i(\mathbf{a}) \cdot \mathbf{v}$$

- So

$$f(\mathbf{a} + \epsilon \mathbf{v}) = f(\mathbf{a}) + \epsilon Df(\mathbf{a}) \cdot \mathbf{v}$$

- $Df(\mathbf{a})$ is **Jacobian matrix** of partial derivatives $\frac{\partial f_i}{\partial x_j}$
- Coefficient of ϵ is Jacobian–vector product

Summary

- We can implement rules to calculate derivatives
- By defining a new type and overloading operations
- Extends directly to derivatives of higher-dimensional functions