

18.330 Problem set 4 (spring 2020)

Submission deadline: 11:59pm on Tuesday, March 3

Exercise 1: Summation two ways

Consider the problem of summing N numbers:

$$S = \sum_{n=1}^N x_n$$

We will compare two different ways to do this with the built-in `sum` function in Julia.

1. Implement the naive summation algorithm as `sum1` that takes in a vector `x` and sums the elements, using e.g. a `for` loop.
2. Implement the summation in a different way (`sum2`) by splitting the vector into two halves and calling the function recursively on each half. The base case when `x` is of length 1 just returns `x[1]`. For simplicity assume that `x` has a length that is a power of two.

The recursive line of your implementation should look something like this:

```
return sum2(view(x, 1:nnew)) + sum2(view(x, nnew+1:n))
```

This is the recursive part. `nnew` is half the length of the input vector, `n/2` (you might want to use `div` to calculate this). While not specifically part of the course, the reason you need to use views is to prevent excessive memory allocations that slow down the code too much for this problem to work – so please use them.

3. Use the built-in Julia `sum` function on `big.(x)` as the true value. Calculate the relative error of `sum1` and `sum2` as the length of `x` is varied from 2^5 to 2^{20} . Make a plot of relative error vs. length. What do you see? Which plot function would you use only given this plot?
4. How do the speeds compare? Using the `@belapsed` function, eg.

```
sum1_time = `@belapsed sum1($x)`
```

benchmark the two `sum` functions that you have defined and plot the run times as a function of length. Which function would you use only given this plot?

5. If everything went according to plan you should have obtained different answers for [3] and [4]. How can we remedy this situation? Recursing all the way down to a base case of length 1 is very inefficient, since there is a time overhead for each function invocation. From your error plot, though, you should be able to see that the errors of your implementations only

start diverging once x reaches a certain length, say 2^{10} (although you can play with this number).

Write a new function `sum3` that uses recursion but that calls `sum1` after x is sufficiently short, i.e. changing the base case specification in your `sum2` function.

[This a technique that we will see repeatedly with other recursive algorithms, such as cache-efficient matrix multiplication and the FFT algorithm.]

6. Calculate the relative error and run time for this new function and also the built-in Julia `sum` function as in [3, 4] and add it the plots you made there. Which function would you use now? Your `sum3` implementation is the same way that Julia calculates sums but with more optimizations.

Hopefully this problem has shown that thinking about errors and how you implement even the simplest algorithms, such as summation, is important if you want to have both performance and accuracy. Naive implementations can be dangerous!

Exercise 2: Conditioning of a problem and stability of an algorithm Consider the problem to calculate

$$\phi(x) = \sqrt{1+x} - 1$$

close to $x = 0$.

1. Calculate the relative condition number $\kappa_\phi(x)$.
2. Is the problem well conditioned near $x = 0$?
3. Consider the obvious algorithm by breaking up the formula into its simplest pieces. Thinking about it as this series of algorithmic steps, calculate the condition number of each step for x near 0. Which is the problematic step?
4. By using an algebraic manipulation, find an alternative algorithm to evaluate the function which is “stable”, i.e. which does not introduce extra numerical error in this way.
5. Implement both the obvious algorithm and your better algorithm. Take the true value of ϕ to be $\phi(\text{big}(x))$ and calculate the relative error of both algorithms for different values of x as it decreases toward 0. Plot the errors as a function of x on the same axes. Is your new algorithm better?

Exercise 3: Conditioning of polynomial roots Consider a degree- n polynomial $p(x) = a_0 + a_1x + \dots + a_nx^n$.

1. Show that the condition number of the problem of finding a root r of this polynomial when the leading coefficient a_n is varied is

$$\kappa = \left| \frac{a_n r^{n-1}}{p'(r)} \right|$$

Hint: To do this use e.g. implicit differentiation on the equation that the root r satisfies.

2. Consider the (in)famous Wilkinson polynomial

$$p(x) = (x - 1)(x - 2) \cdots (x - 20).$$

Calculate the coefficients of the polynomial using e.g. your `Polynomial` type from a previous problem set or the `Polynomials.jl` package. Note that you will need to use `Int128` or `BigInt`.

3. Calculate, e.g. analytically or using automatic differentiation, the condition number of each root of the polynomial. Which roots are well-conditioned and which are ill-conditioned?
4. Use e.g. the `roots` function from the `PolynomialRoots.jl` package to calculate the roots of $p(x)$. Are they correct?

Now perturb a single coefficient of $p(x)$ by `'randn()*2.0^(-23)'` and find and plot the roots on the same graph. Repeat this 50 times to see a visual representation of the unstable nature of polynomial root finding.

Does the result agree with what you calculated in [3]?

Exercise 4: Lagrange interpolation

1. Given nodes t_i , write a function to calculate the Lagrange cardinal functions $\ell_k(x)$ that satisfies $\ell_k(t_i) = [i = k]$ (remember $[i = k]$ is 0 unless $i = k$). Your function should take in a vector `t`, a value `x` and the index `k` and return the value of $\ell_k(x)$. What is the operation count for calculating the cardinal function?
2. Write a function to calculate the Lagrange interpolant in nodes t_i with data y_i at the point x using your lagrange cardinal function from [1]. What is the operation count for calculating the Lagrange interpolant?
3. Find the Lagrange interpolant to the function $\exp(x)$ sampled at $N + 1$ equally spaced points in the interval $[-1, 1]$. Make a plot of the interpolant over this interval when $N = 10$. Also add a plot of $\exp(x)$ using the built in function and a scatter plot of the interpolating points. Do you see what you expect.

4. Calculate numerically the error in the interpolant compared to the function $\exp(x)$ by sampling at, say, 100 points in the interval and taking the maximum error. How does the error vary with the degree of the interpolant?
5. Consider the function $f(x) = \frac{1}{1+25x^2}$ on the interval $[-1, 1]$. What happens as you increase the number of equally-spaced points? Make an interactive visualization. This is known as the **Runge phenomenon**
6. Instead, interpolate at the **Chebyshev points**, $t_i := \cos(\theta_i)$, where θ_i are equally-spaced angles between 0 and π . Make another interactive visualization. You should find that now we do have convergence. Find numerically how the rate of convergence of the distance between the interpolant and the true function behaves as the number of points increases. This is known as **spectral convergence**.